

Analysis of Trade-offs in Fault-Tolerant Distributed Computing and Replicated Databases

Anatoliy Gorbenko^{1,2}

¹Leeds Beckett University

Leeds, UK

A.Gorbenko@leedsbeckett.ac.uk

Andrii Karpenko

²National Aerospace University

Kharkiv, Ukraine

A.Karpenko@student.csn.khai.edu

Olga Tarasyuk

²National Aerospace University

Kharkiv, Ukraine

O.Tarasyuk@csn.khai.edu

Abstract—This paper examines fundamental trade-offs in fault-tolerant distributed systems and replicated databases built over the Internet. We discuss interplays between consistency, availability, and latency which are in the very nature of globally distributed computer systems and also analyse their interconnection with durability and energy efficiency. In this paper we put forward an idea that consistency, availability, latency, durability and other properties need to be viewed as more continuous than binary in contrast to the well-known CAP/PACELC theorems. We compare different consistency models and highlight the role of the application timeout, replication factor and other settings that essentially determine the interplay between above properties. Our findings may be of interest to software engineers and system architects who develop Internet-scale distributed computer systems and cloud solutions.

Keywords—distributed system, replication, trade-offs, consistency, availability, latency, durability, energy-efficiency

I. INTRODUCTION

Internet-scale distributed computer systems are now extensively used in business-critical applications and critical infrastructures. In such application domains, system failures affect businesses and people's lives. Thus, ensuring dependability of distributed computer systems is a must, as well as a challenge. However, by their nature, large-scale distributed systems running over the Internet are subject to components failures, packets loss, network disconnections, congestions and other accidents.

High availability requirements for many Internet applications call for the use of data replication and system redundancy. Traditional fault tolerance mechanisms such as N-modular, cold- and hot-spare redundancy usually rely on a synchronous communication between replicated components. This suggests that system replicas are synchronized over a short and well-predicted amount of time [1]. This is a reasonable assumption for embedded applications and for those distributed computer systems which components are located, for instance, within the same data center or in the same local area network. However, this does not apply to globally distributed computer systems, which replicas are deployed across the Internet and their updates cannot be propagated immediately. This circumstance makes it difficult to guarantee consistency across replicas.

The Internet and globally distributed computer systems are characterized by a high level of uncertainty of network delays. This makes it almost impossible to guarantee that network messages will always be delivered between system components within a certain time. It has been previously shown that there is a significant uncertainty of response time in service-oriented systems invoked over clouds and the

Internet [2]. Besides, in our practical experiments and as discussed by other researchers [3, 4, 5], failures occur regularly on the Internet, clouds and in scale-out data center networks.

When system architects employ replication and other fault tolerant techniques for the Internet- and cloud-based systems, they have to care about additional delays and their uncertainty and also understand energy and other overheads. Besides, maintaining consistency between replicas is another important issue that needs to be addressed in fault-tolerant distributed computing and replicated data storages.

Maintaining several system replicas inevitably increases the overall energy, consumed by the system. The bigger the replication factor, the higher availability can be ensured at higher energy costs. Moreover, necessity of data synchronization between replicas to guaranty their consistency causes additional energy overheads spent for parallel requests processing and transferring larger amount of data over the network to propagate updates.

This paper examines fundamental trade-offs between system latency, durability and energy consumption in addition to Consistency (C), Availability (A) and Partition tolerance (P) properties, as described by the CAP theorem [6]. In this work we put forward an idea that these properties need to be viewed as more continuous than binary. Understanding trade-offs between them will allow systems developers to predict system response time depending on the used replication factor and/or the selected consistency level. Besides, it will enable balancing availability, durability and/or consistency against latency, power consumption and other QoS properties.

The rest of the paper is organized as follows. In Sections 2 we discuss fundamental CAP and PACELC theorems and their qualitative implications. Section 3 discusses different consistency models and examines trade-offs in fault-tolerant distributed computing and replicated databases between CAP properties, latency, durability and energy consumption. Finally, we draw our conclusions in Section 4 where the role of the application timeout, replication factor and other system settings that essentially determine the interplay between above properties is also discussed.

II. CAP AND PACELC THEOREMS AND THEIR IMPLICATIONS

The CAP theorem [6], first appeared in 1998-1999, defines a qualitative trade-off between system Consistency, Availability, and Partition tolerance. It declares that the only two of the three properties can be preserved at once in distributed replicated systems.

Gilbert and Lynch [7] consider the CAP theorem as a particular case of a more general trade-off between consistency

and availability in unreliable distributed systems propagating updates eventually over time.

Indeed, Internet-scale distributed systems cannot avoid partitions happened due to network failures and congestions, arbitrary message losses and components crashes and, hence, have to choose between strong consistency or high availability. Failing to achieve consistency (i.e. receive responses from all replicas) within the specified timeout causes a partition of the replicated systems. Thus, a distributed system working over the unreliable and unpredictable network environment has to sacrifice one of these two properties.

Timeout settings are of great importance in the context of the CAP theorem. If timeout is less than the typical response time, a system will likely enter a partition mode more often [8]. When a system detects a partition (e.g. when one of its replica did not respond before the time-out) it has to decide whether to return a possibly stale result to a client or to reply with an exception message notifying about service unavailability. Thus, the CAP theorem suggests the following three types of systems:

- CA, e.g. traditional ACID-oriented RDBMS systems (MySQL, MS SQL, Oracle, PostgreSQL, etc.) which preserve Consistency and Availability;
- AP, NoSQL databases such as Cassandra, Riak, CouchDB, Voldemort, Dynamo which relax Consistency in favor of Availability and Partition tolerance;
- CP, NoSQL databases such as HBase, MongoDB, Redis, BigTable, MemcacheDB, which preserve Consistency when Partitioned.

The PACELC [9] theorem is a further refinement of CAP. It suggests that that in case of *partitioning* (P) of a distributed computer system, one has to choose between *availability* (A) and *consistency* (C), *else* (E) in the absence of partitions the replicated systems still face a trade-off between *latency* (L) and *consistency* (C). The PACELC theorem defines the following types of distributed replicated systems:

- PC/EC, e.g. BigTable, HBase, VoltDB/H-Store, Megastore;
- PC/EL, e.g. PNUTS;
- PA/EL, e.g. Dynamo, Cassandra, and Riak;
- PA/EC, e.g. MongoDB.

However, a notion of *Partition* is not well defined in the PACELC theorem as well. For instance, PC (PA) does not indicate that the system is fully consistent (available). It should rather be interpreted as ‘*if partitioning happens it causes more consistency (availability) issues than availability (consistency) issues*’.

Though CAP and PACELC theorem help developers to understand system trade-offs between consistency and availability/latency, there are no methods available that allow interplaying consistency against availability and latency in a quantitative way. Besides, the CAP and PACELC theorems do not take into account other fundamental trade-offs, e.g. between durability and latency, energy consumption, availability and consistency, etc.

III. TRADE-OFFS IN FAULT-TOLERANT DISTRIBUTED COMPUTING AND REPLICATED DATABASES

In this paper we put forward an idea that CAP/PACELC properties need to be viewed as more continuous than binary.

Indeed, *availability* is measured as usual between 0% and 100%. *Latency* (response time) can theoretically vary between zero and infinity. Though, in practice it is restricted from the right by the application timeout and from the left by some minimal response time higher than zero. Consequently, the replica’s timeout defines system *partitioning*.

Consistency is also a continuum, ranging from weak consistency at one extreme to strong consistency on the other, with varying points of relaxed (e.g. eventual) consistency levels in between.

In this section we describe different consistency models and discuss trade-offs between core QoS properties of distributed computer systems.

A. Consistency Models and Levels

Data consistency models define a contract between a replicated data store and its users, in which the data store specifies guarantees on the results of read and write operations in the presence of concurrent users’ requests.

There are two major groups of consistency models [10]. The first one guarantees the certain state of a data store for all users (data-centric models); the second one provides guarantees only for an individual user (client-centric models) while the data seen can be varying from users to users. There is a variety of different consistency models used in distributed computing and storage systems [10] which are structured in Fig. 1.

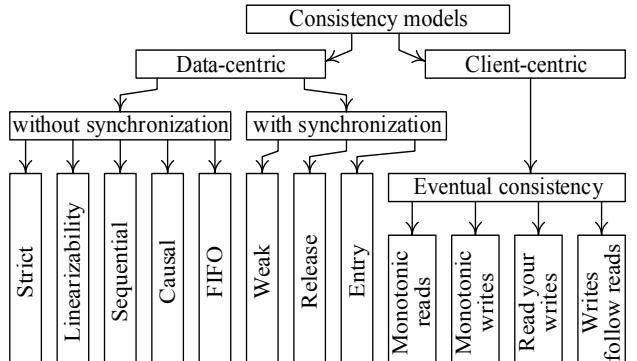


Fig. 1. A hierarchy of consistency models

The strong consistency cannot be efficiently achieved in replicated systems. Thus, many distributed database systems implements different kinds of relaxed/eventual consistency models.

Unfortunately, there are no general rules for relaxing consistency. As a result, different vendors implement different consistency models, which are hardly compatible and are difficult to match with each other.

For example, *Apache Cassandra* defines a discrete set of consistency levels specifying a number of replicas queried in each read and write operation (see Table I). The strong consistency guaranteeing that a read will always reflect the most recent write is achieved when a sum of nodes written and nodes read is higher than data replication factor.

TABLE I. APACHE CASSANDRA CONSISTENCY MODEL

Consistency level	Definition and consistency guarantee
ONE	Data must be written to the commit log and <i>memtable</i> of at least one replica node before acknowledging the write operations to a client; when reading data, Cassandra queries and returns a response from a single replica (the nearest replica with the least network latency); the strong consistency is guaranteed when a sum of nodes written and nodes read is higher than data replication factor.
TWO	Data must be written to at least two replica nodes before being acknowledged; read operations will return the most recent record from two of the closest replicas (the most recent data is determined by comparing timestamps of records returned by those two replica)
THREE	Similar to TWO but for three replicas
QUORUM	A quorum of nodes needs to acknowledge the write or to return a response for a read request; a quorum is calculated by rounding down to a whole number the following estimate: <i>replication factor/2+1</i>
ALL	Data must be written to all replica nodes in a cluster before being acknowledged; read requests return the most recent record after all replicas have responded. The read operation will fail even if a single replica does not respond
EACH_QUORUM, LOCAL_QUORUM, LOCAL_ONE	Additional consistency levels which become available if Cassandra runs across multiple data centres

TABLE II. MONGODB CONSISTENCY MODEL

Consistency level	Definition and consistency guarantee
w:0, j:false	The weakest consistency setting (writes can be lost even without partition) which provides the lowest latency
w:1, j:false	Writes are guaranteed onto disk of the primary replica; this provides very low latency but very weak consistency
w:2, j:false	Writes are guaranteed on the primary replica's disk and in the memory of one of the secondary replicas; this provides low latency and low consistency
w:2, j: true	Writes are guaranteed on the disks of primary replica and one of the secondary replicas; this provides medium latency and consistency
w:majority, j: false	Writes are guaranteed on the primary replica's disk and in the memory of a majority of secondary replicas; this provides medium latency and consistency
w:majority, j: true	Writes are guaranteed on the disks of primary replica and a majority of secondary replicas; this provides high latency and consistency

TABLE III. AZURE COSMOS DB CONSISTENCY MODEL

Consistency level	Definition and consistency guarantee
STRONG (Reads: local minority; Writes: global majority)	Strong consistency offers a linearizability guarantee, e.g. the reads are guaranteed to return the most recent committed writes with a zero staleness window.
BOUNDED STALENESS (Reads: local minority; Writes: local majority)	It is guaranteed that reads never see out-of-order writes. Though, reads might lag behind writes by at most K updates of a record or by T time interval (i.e. a staleness window) whichever is less.
SESSION (Reads: single replica with session token; Writes: local majority)	Within a single client session it is guaranteed that reads never see out-of-order writes; monotonic reads, writes, write-follows-reads and read-your-writes are also guaranteed; For other client sessions
CONSISTENT PREFIX (Reads: single replica; Writes: local majority)	It is guaranteed that read never see out-of-order writes or writes with gaps; e.g. it is guaranteed to observe an ordered sequence of writes (starting with the first one) that stored at the master replica at some time in the past; more recent writes can be missed.
EVENTUAL (Reads: Single replica; Writes: local majority)	There is no ordering guarantee for reads. In the absence of any further writes, the replicas eventually converge; users may read the values that are older than the ones it had read before.

MongoDB's consistency model is based on tuning w (write concern) and j (journaling) parameters as shown in Table II. Writes are always replicated asynchronously from primary to secondary replicas. By default, reads are done on primary replica. Enabling reads from secondaries makes *MongoDB* always eventually consistent.

Azure Cosmos DB is a cloud-based storage system which can be globally distributed across multiple Azure regions. It can be configured for a single- (i.e. single write region; by default) or multi-master replication and supports five consistency levels described in Table III.

B. Trade-offs Between Consistency, Availability and Latency

System consistency, availability and latency are tightly connected. Availability can be interpreted as a probability that each request eventually receives a response. Though, in many real systems a response that is too late (i.e. beyond the application timeout) is treated as a failure.

However, high latency has an undesirable effect for many interactive web applications. Experience report [11] shows that the probability that a client will continue to use an on-line

system or a website is significantly reduced if the response time is increased even as small as 100 ms.

A system or its replica could be considered as unavailable if the actual response time exceeds the application time out. I.e. a partition can be considered as a time bound on replica's response time. Thus, in term of CAP a slow network connection or slowly responding replica (for example, due to high requests load) may lead to a decision that the system is partitioned. Nowadays, architects of distributed database management systems and large-scale web applications like Twitter, Instagram, Facebook, etc. often prefer to weak consistency guarantee by introducing asynchronous data updates in favour of high system availability and low response time. However, the most promising approach is to try to balance these properties with regards to the desired latency and required consistency.

Our interpretation of the CAP and PACELC theorems and the trade-offs resulting from them is depicted in Fig. 2. The consistency model we use determines the number of replicas which are queried simultaneously to return the adjudicated (consistent) result to a client application (this is pretty similar to Cassandra consistency model; see Table I).

In particular, the following consistency levels are possible:

- ONE (equivalent to a hot-spare redundancy); A response from a single replica is forwarded to the client. This is the weakest consistency level though it guarantees the minimal latency. We can further consider a few variations of this consistency level:
 - ONE-ONE – a single replica (usually the nearest replica with the least network latency) is queried whose response is returned to a user;
 - ONE-ALL – all replicas are queried simultaneously and the fastest response is returned to a user without waiting for other replicas' responses, which are ignored;
 - ONE-QUORUM – a quorum of replicas is requested and the fastest response is returned to a user;
- ALL (equivalent to a N-modular redundancy) – the system must wait until ALL replicas return their responses. In this case the response time is constrained by the slowest replica though the strongest consistency is provided.
- QUORUM (equivalent to two-out-of-three or majority systems) – the system must wait for the responses from a QUORUM of replica web services. It provides a compromise between ONE and ALL levels trading off latency versus consistency.

If minimum latency is the top priority users should consider a weaker consistency level, e.g. ONE and its variations. ONE-ALL level could provide faster response (faster than ONE-ONE and ONE-QUORUM) at higher computational costs and data transfer overheads.

If consistency is the top priority, users should opt for a stronger consistency setting, which, however, worsens system latency. The system always ensures the strong consistency guaranteeing that a read will always reflect the most recent write by preserving the following rule:

$$(nodes_written + nodes_read) > replication_factor \quad (1)$$

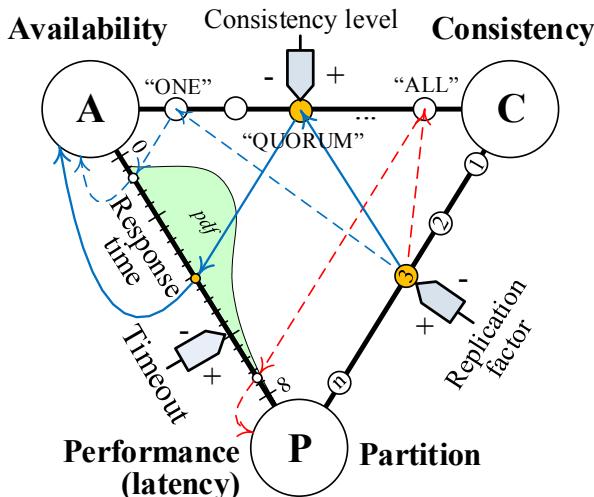


Fig. 2. The CAP/PACELC trade-offs

For example, QUORUM consistency level used for both write and read operations always ensures the strong consistency, trading off between reads and writes latencies. The strong read consistency is also provided when the ALL consistency level is used to read data while the ONE consistency level is used to write and vice versa. In the first case the preference is given to minimizing writes latency, while in the second case the minimum reads latency is ensured.

Application timeout is considered as a bound between system availability and performance (in term of latency/response time) [12]. Besides, if replica does not respond until the specified timeout, the system is considered as partitioned. Thus, system designers should be able to set up timeouts according to the desired system response time also keeping in mind a choice between consistency and availability.

C. Trade-offs Between Performance, Consistency and Durability

Durability is the ability of a system to keep the committed data consistent after crashes, drive failures, power outages, or other forms of corruption. Storing data in memory reduces querying time and provides more predictable and faster performance than storing data on a hard or solid state drive.

Many NoSQL databases store data in memory which improves performance giving up durability. Even traditional RDBMS like PostgreSQL, Oracle, MS SQL or MySQL can be configured to enable periodic durable commits or in-memory data storage mode.

Durability of such systems is ensured by *transaction/commit logging*, which records changes to the database in a journal file before they are committed to a user. It is used for in-memory data recovery in case of power loss or rollback operation. However, a database can be configured to store the commit/transaction log in memory and flush it to disk only periodically. This improves performance however creates so called 'data loss window' [13].

For example, Cassandra NoSQL flushes a commit log to disk every 10 seconds by default. Thus, if a power outage happens right before writing the commit log to disk a system could potentially lose up to ten seconds of data.

Nevertheless, *transaction/commit logging* can not prevent data loss in the event of a disk failure. If this happens, lost data could be restored on the crashed node (when it is being repaired) only if the system replicates data across multiple nodes (e.g. a replication factor is above one).

Thus, a higher replication factor improves durability property of a system. However, it affects system consistency and latency depending on the chosen consistency level, as it is shown in Fig. 3.

D. Trade-offs Between CAP properties, Fault-Tolerance and Energy Consumption

Modern ICT industry, according to Gartner's research [14], is responsible for approximately 2 percent of global CO₂ emissions. This is equivalent to the CO₂ emissions of the whole aviation sector of economy.

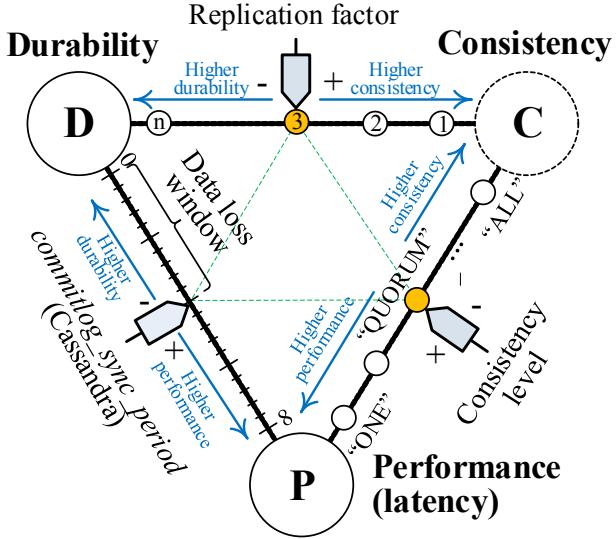


Fig. 3. Trade-offs between Durability, Performance and Consistency

Alex Wissner-Gross reported [15] that a typical Google search request generates about 7g of CO₂ (though Google argues that it is only about 0.2 g). This is half as much as boiling a kettle. Besides, surfing the Internet generates approximately 20 milligrams of CO₂ per second on average.

The total energy consumption of all computing and communication equipment in the world accounts for 160 GW per year, which is about 8% of the total world's energy generation [16]. It is also worth noting that only 1 of the 27 Watts consumed by a typical data center is spent directly to perform user computations [17]. Thus, the total aggregated energy loss in data centers can reach up to 97%.

These and other facts reveal the significance of ICT contribution into the word's energy consumption and CO₂ emission. Thus, enhancing energy efficiency of hardware, software and communication technologies is one of the key issues of the modern ICT industry.

Fig. 4 presents trade-offs between CAP properties, Fault-Tolerance (FT) and Energy Consumption (EC). Replication factor defines the main trade-off between fault tolerance and energy consumption which is proportional to the number of replicas.

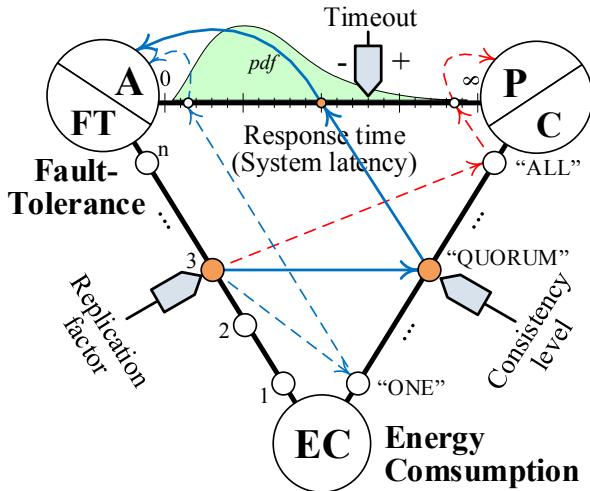


Fig. 4. Trade-offs between Energy Consumption and CAP properties

Developers of distributed computer systems employ replication (i.e. redundancy) to pursue two major goals. Firstly, it improves availability and durability by tolerating crashes, errors and failures occurred in such systems and in the communication environment. Global replica distribution across different availability zones in the Internet also guarantees disaster tolerance and survivability. Secondly, replication enhances system performance by balancing user requests between replicated nodes.

At the same time, a high degree of redundancy (i.e. a large replication factor), which implies better fault-tolerance, does not always guarantee high availability. Indeed, using a high replication factor together with strong consistency settings only increases system latency and reduces the probability of a system to respond before the specified time-out (see Fig. 4). System latency and its uncertainty further increase when replicas are distributed across the Internet. Hosting all replicas in the same data center, in general, reduces a deviation between their response times; though, the probability of a common-mode failure is increased. Besides, users from different geographic locations will experience significant differences in system latency and its variation.

Another important aspect of the used consistency level is the number of replicas that are invoked simultaneously to execute a particular read or write request. The higher consistency level the more replicas are requested simultaneously. This increases the overall energy consumption taking into account additional power spent on transmitting requests/responses and propagating updates over the network. In turn, the greater distance between replicas, the more energy is consumed.

Thus the replication factor and consistency level contribute together to the overall energy consumption. The replication factor can be seen as the dominant aspect of system energy consumption, while consistency level adds a variable component.

IV. CONCLUSION

When employing replication and other fault-tolerance techniques over the Internet and clouds, engineers have to deal with delays, their uncertainty, timeouts, adjudication of asynchronous replies from replicas, and other specific issues involved in global distributed systems. The overall aim of this work is studying fundamental trade-offs in distributed database systems and fault tolerant Internet computing and also examining interplays between CAP properties, latency, durability and energy consumption.

Distributed nature of modern computer applications increases a probability of failures. Replication helps to ensure system fault tolerance and to increase its performance by load balancing. At the same time, running several replicas proportionally increases energy consumption. Besides, a replicated system causes the consistency issue. Necessity to provide strong consistency requires concurrent invocation of several replicas that additionally increases the overall energy consumption. Consequently, stronger data consistency worsens system latency. This finding confirms one of the generally adopted qualitative implications of the CAP theorem [6].

In the paper we reveal key system settings that can be used to interplay CAP properties, latency, durability and energy consumption. They include:

- *Replication factor* (i.e. a number of system replicas).
- *Consistency level* which defines consistency guarantees which a system provides to users).
- *Time-out settings* (i.e. how long a client should wait for read/write operations to complete and how long the system should wait for replicas' responses before it is considered partitioned).
- *Commit log synchronization time* which defines how often the commit log buffer stored in memory is synchronized to disk.

One should note that these settings are tightly connected and should not be considered individually. For instance, a replication factor higher than one always causes consistency issue. However, depending on the consistency settings higher replication factor can further increase system latency (if a system is configured to provide more strict consistency guarantees) or reduce it (if consistency is more relaxed).

Though existing researches help to define a series of useful qualitative implications of the CAP/PACELC theorems (e.g. '*better consistency worsens system availability and latency*'), developers have not provided yet with adequate quantitative models helping to predict system response time corresponding to the chosen consistency level and to precisely trade-off between other CAP properties. Estimation of the system worst-case execution time still remains a common practice for many applications (e.g. embedded computer systems, server fault-tolerance solutions, like STRATUS, etc.). However this approach is no longer a viable solution for the large-scale computer systems which replicas/components are globally distributed across the Internet.

In our previous works (e.g. [2, 5]) we demonstrated that unpredictable extreme delays exceeding the value of ten average response times could happen in such system quite often. Thus, system architects need novel analytical models providing a quantitative basis for the system response time prediction depending on the consistency level provided to (or requested by) clients.

REFERENCES

- [1] P. Lee and T. Anderson, Fault Tolerance. Principles and Practice, Wien - New-York: Springer-Verlag, 1990, p. 320.
- [2] A. Gorbenko, A. Romanovsky, O. Tarasyuk, V. Kharchenko and S. Mamutov, "Exploring Uncertainty of Delays as a Factor in End-to-End Cloud Response Time," in 9th European Dependable Computing Conference (EDCC'2012), Sibiu, Romania, 2012.

- [3] R. Potharaju and N. Jain, "When the Network Crumbles: An Empirical Study of Cloud Network Failures and their Impact on Services," in 4th ACM Symposium on Cloud Computing (SOCC'2013), Santa Clara, CA, 2013.
- [4] C. Scott, D. Choffnes, I. Cunha and e. al, "LIFEGUARD: practical repair of persistent route failures," in ACM SIGCOMM 2012 conference on Applications, technologies, architectures, and protocols for computer communication, New York, USA, 2012.
- [5] A. Gorbenko, V. Kharchenko, O. Tarasyuk, Y. Chen and A. Romanovsky, "The threat of uncertainty in Service-Oriented Architecture," in RISE/EFTS Joint International Workshop on Software Engineering for Resilient Systems, SERENE'08, Newcastle, 2008.
- [6] E. Brewer, "Towards Robust Distributed Systems," in 19th Annual ACM Symposium on Principles of Distributed Computing, Portland, USA, 2000.
- [7] S. Gilbert and N. Lynch, "Brewer's Conjecture and the Feasibility of Consistent, Available, Partition-Tolerant Web Services," ACM SIGACT News, vol. 33, no. 2, pp. 51-59, 2002.
- [8] A. Gorbenko, A. Romanovsky and O. Tarasyuk, "Fault tolerant internet computing: Benchmarking and modelling trade-offs between availability, latency and consistency," Journal of Network and Computer Applications, vol. 146, p. 102412, 2019.
- [9] D. Abadi, "Consistency Tradeoffs in Modern Distributed Database System Design," IEEE Computer, vol. 45, no. 2, pp. 37-42, 2012.
- [10] A. Tanenbaum and M. Van Steen, Distributed systems: Principles and Paradigms, Pearson Prentice Hall, 2006, p. 704.
- [11] J. Brutlag, "Speed Matters for Google Web Search," Google, Inc., 22 June 2009. [Online]. Available: http://services.google.com/fh/files/blogs/google_delayexp.pdf. [Accessed 01 07 2019].
- [12] A. Gorbenko and A. Romanovsky, "Time-outing Internet Services," IEEE Security & Privacy, vol. 11, no. 2, pp. 68-71, 2013.
- [13] R. Alagappan, A. Ganesan, Y. Patel, T. S. Pillai, A. C. Arpacı-Dusseau and R. H. Arpacı-Dusseau, "Correlated Crash Vulnerabilities," in 12th USENIX Symposium on Operating Systems Design and Implementation, Savannah, USA, 2016.
- [14] Gartner, Inc., "Gartner Estimates ICT Industry Accounts for 2 Percent of Global CO2 Emissions," 2007. [Online]. Available: <http://www.gartner.com/newsroom/id/503867>.
- [15] J. Leake and R. Woods, "Revealed: the environmental impact of Google searches," 11 January 2009. [Online]. Available: <https://www.thetimes.co.uk/article/revealed-the-environmental-impact-of-google-searches-xvcc72q8t2z>.
- [16] T. Jorg, "Efficiency starts with the power supply," Journal of network solutions/LAN, vol. 4, 2013.
- [17] D. Neilson, "IBM and Dynamic Infrastructure," IBM Systems Group, March 2009. [Online]. Available: www.nesc.ac.uk/talks/968/NESC_Neilson.ppt.